

# Mini-Project

Brittney Hayes

## 1. Preparing the data

```
#View(WisconsinCancer)

# Save input data file into your Project directory
fna.data <- "WisconsinCancer.csv"

# Complete the following code to input the data and store as wisc.df
wisc.df <- read.csv(fna.data, row.names = 1)

# View the first few rows of the dataframe
head(wisc.df)
```

	X	X.1	X.2	X.3	X.4
id	diagnosis	radius_mean	texture_mean	perimeter_mean	area_mean
842302	M	17.99	10.38	122.8	1001
842517	M	20.57	17.77	132.9	1326
84300903	M	19.69	21.25	130	1203
84348301	M	11.42	20.38	77.58	386.1
84358402	M	20.29	14.34	135.1	1297

	X.5	X.6	X.7	X.8
id	smoothness_mean	compactness_mean	concavity_mean	concave points_mean
842302	0.1184	0.2776	0.3001	0.1471
842517	0.08474	0.07864	0.0869	0.07017
84300903	0.1096	0.1599	0.1974	0.1279
84348301	0.1425	0.2839	0.2414	0.1052
84358402	0.1003	0.1328	0.198	0.1043

	X.9	X.10	X.11	X.12	X.13
id	symmetry_mean	fractal_dimension_mean	radius_se	texture_se	perimeter_se
842302	0.2419	0.07871	1.095	0.9053	8.589

842517	0.1812		0.05667	0.5435	0.7339	3.398
84300903	0.2069		0.05999	0.7456	0.7869	4.585
84348301	0.2597		0.09744	0.4956	1.156	3.445
84358402	0.1809		0.05883	0.7572	0.7813	5.438

	X.14	X.15	X.16	X.17	X.18
id	area_se	smoothness_se	compactness_se	concavity_se	concave points_se
842302	153.4	0.006399	0.04904	0.05373	0.01587
842517	74.08	0.005225	0.01308	0.0186	0.0134
84300903	94.03	0.00615	0.04006	0.03832	0.02058
84348301	27.23	0.00911	0.07458	0.05661	0.01867
84358402	94.44	0.01149	0.02461	0.05688	0.01885

	X.19	X.20	X.21	X.22
id	symmetry_se	fractal_dimension_se	radius_worst	texture_worst
842302	0.03003	0.006193	25.38	17.33
842517	0.01389	0.003532	24.99	23.41
84300903	0.0225	0.004571	23.57	25.53
84348301	0.05963	0.009208	14.91	26.5
84358402	0.01756	0.005115	22.54	16.67

	X.23	X.24	X.25	X.26
id	perimeter_worst	area_worst	smoothness_worst	compactness_worst
842302	184.6	2019	0.1622	0.6656
842517	158.8	1956	0.1238	0.1866
84300903	152.5	1709	0.1444	0.4245
84348301	98.87	567.7	0.2098	0.8663
84358402	152.2	1575	0.1374	0.205

	X.27	X.28	X.29
id	concavity_worst	concave points_worst	symmetry_worst
842302	0.7119	0.2654	0.4601
842517	0.2416	0.186	0.275
84300903	0.4504	0.243	0.3613
84348301	0.6869	0.2575	0.6638
84358402	0.4	0.1625	0.2364

	X.30
id	fractal_dimension_worst
842302	0.1189
842517	0.08902
84300903	0.08758
84348301	0.173
84358402	0.07678

```
# We can use -1 here to remove the first column
wisc.data <- wisc.df[,-1]

# Create diagnosis vector for later
diagnosis <- wisc.df$X
```

## 1. Exploratory Data Analysis

Q1. How many observations are in this dataset?

```
num_observations <- nrow(wisc.data)
```

Answer: 570

Q2. How many of the observations have a malignant diagnosis?

```
num_malignant <- sum(diagnosis == "M")
```

Answer: 212

Q3. How many variables/features in the data are suffixed with `_mean`?

```
num_mean_variables <- sum(grep("_mean", names(wisc.data)))
```

Answer: 0

## 2. Principal Component Analysis

```
str(wisc.data)
```

```
'data.frame':  570 obs. of  30 variables:
 $ X.1 : chr  "radius_mean" "17.99" "20.57" "19.69" ...
 $ X.2 : chr  "texture_mean" "10.38" "17.77" "21.25" ...
 $ X.3 : chr  "perimeter_mean" "122.8" "132.9" "130" ...
 $ X.4 : chr  "area_mean" "1001" "1326" "1203" ...
 $ X.5 : chr  "smoothness_mean" "0.1184" "0.08474" "0.1096" ...
 $ X.6 : chr  "compactness_mean" "0.2776" "0.07864" "0.1599" ...
 $ X.7 : chr  "concavity_mean" "0.3001" "0.0869" "0.1974" ...
 $ X.8 : chr  "concave points_mean" "0.1471" "0.07017" "0.1279" ...
```

```

$ X.9 : chr "symmetry_mean" "0.2419" "0.1812" "0.2069" ...
$ X.10: chr "fractal_dimension_mean" "0.07871" "0.05667" "0.05999" ...
$ X.11: chr "radius_se" "1.095" "0.5435" "0.7456" ...
$ X.12: chr "texture_se" "0.9053" "0.7339" "0.7869" ...
$ X.13: chr "perimeter_se" "8.589" "3.398" "4.585" ...
$ X.14: chr "area_se" "153.4" "74.08" "94.03" ...
$ X.15: chr "smoothness_se" "0.006399" "0.005225" "0.00615" ...
$ X.16: chr "compactness_se" "0.04904" "0.01308" "0.04006" ...
$ X.17: chr "concavity_se" "0.05373" "0.0186" "0.03832" ...
$ X.18: chr "concave points_se" "0.01587" "0.0134" "0.02058" ...
$ X.19: chr "symmetry_se" "0.03003" "0.01389" "0.0225" ...
$ X.20: chr "fractal_dimension_se" "0.006193" "0.003532" "0.004571" ...
$ X.21: chr "radius_worst" "25.38" "24.99" "23.57" ...
$ X.22: chr "texture_worst" "17.33" "23.41" "25.53" ...
$ X.23: chr "perimeter_worst" "184.6" "158.8" "152.5" ...
$ X.24: chr "area_worst" "2019" "1956" "1709" ...
$ X.25: chr "smoothness_worst" "0.1622" "0.1238" "0.1444" ...
$ X.26: chr "compactness_worst" "0.6656" "0.1866" "0.4245" ...
$ X.27: chr "concavity_worst" "0.7119" "0.2416" "0.4504" ...
$ X.28: chr "concave points_worst" "0.2654" "0.186" "0.243" ...
$ X.29: chr "symmetry_worst" "0.4601" "0.275" "0.3613" ...
$ X.30: chr "fractal_dimension_worst" "0.1189" "0.08902" "0.08758" ...

```

```

# Had issues "Error in colMeans(wisc.data) : 'x' must be numeric" Converting to numeric (i
wisc.data$X.1<- as.numeric(as.factor(wisc.data$X.1))
wisc.data$X.2<- as.numeric(as.factor(wisc.data$X.2))
wisc.data$X.3<- as.numeric(as.factor(wisc.data$X.3))
wisc.data$X.4<- as.numeric(as.factor(wisc.data$X.4))
wisc.data$X.5<- as.numeric(as.factor(wisc.data$X.5))
wisc.data$X.6<- as.numeric(as.factor(wisc.data$X.6))
wisc.data$X.7<- as.numeric(as.factor(wisc.data$X.7))
wisc.data$X.8<- as.numeric(as.factor(wisc.data$X.8))
wisc.data$X.9<- as.numeric(as.factor(wisc.data$X.9))
wisc.data$X.10<- as.numeric(as.factor(wisc.data$X.10))
wisc.data$X.11<- as.numeric(as.factor(wisc.data$X.11))
wisc.data$X.12<- as.numeric(as.factor(wisc.data$X.12))
wisc.data$X.13<- as.numeric(as.factor(wisc.data$X.13))
wisc.data$X.14<- as.numeric(as.factor(wisc.data$X.14))
wisc.data$X.15<- as.numeric(as.factor(wisc.data$X.15))
wisc.data$X.16<- as.numeric(as.factor(wisc.data$X.16))
wisc.data$X.17<- as.numeric(as.factor(wisc.data$X.17))
wisc.data$X.18<- as.numeric(as.factor(wisc.data$X.18))

```

```

wisc.data$X.19<- as.numeric(as.factor(wisc.data$X.19))
wisc.data$X.20<- as.numeric(as.factor(wisc.data$X.20))
wisc.data$X.21<- as.numeric(as.factor(wisc.data$X.21))
wisc.data$X.22<- as.numeric(as.factor(wisc.data$X.22))
wisc.data$X.23<- as.numeric(as.factor(wisc.data$X.23))
wisc.data$X.24<- as.numeric(as.factor(wisc.data$X.24))
wisc.data$X.25<- as.numeric(as.factor(wisc.data$X.25))
wisc.data$X.26<- as.numeric(as.factor(wisc.data$X.26))
wisc.data$X.27<- as.numeric(as.factor(wisc.data$X.27))
wisc.data$X.28<- as.numeric(as.factor(wisc.data$X.28))
wisc.data$X.29<- as.numeric(as.factor(wisc.data$X.29))
wisc.data$X.30<- as.numeric(as.factor(wisc.data$X.30))

# Check column means and standard deviations
apply(wisc.data,2,sd)

```

	X.1	X.2	X.3	X.4	X.5	X.6	X.7	X.8
129.7166	133.6683	151.6142	155.2636	136.4445	154.3905	158.5430	159.7768	
	X.9	X.10	X.11	X.12	X.13	X.14	X.15	X.16
118.4978	141.1332	155.3676	145.5543	152.5752	151.1658	157.6017	154.7020	
	X.17	X.18	X.19	X.20	X.21	X.22	X.23	X.24
155.7535	146.3378	140.5687	156.2770	127.3364	145.1813	150.9367	157.8110	
	X.25	X.26	X.27	X.28	X.29	X.30		
109.8857	151.2856	157.8129	144.4117	142.2252	154.0813			

```

numeric_data <- wisc.data[, c(1:30)]
means <- colMeans(numeric_data)

# Perform PCA on wisc.data (now numeric_data)
wisc.pr <- prcomp(numeric_data, scale. = TRUE)

# Look at summary of results
summary(wisc.pr)

```

Importance of components:

	PC1	PC2	PC3	PC4	PC5	PC6	PC7
Standard deviation	3.4186	1.8552	1.72721	1.40694	1.31863	1.16598	1.06146
Proportion of Variance	0.3896	0.1147	0.09944	0.06598	0.05796	0.04532	0.03756
Cumulative Proportion	0.3896	0.5043	0.60372	0.66971	0.72767	0.77298	0.81054
	PC8	PC9	PC10	PC11	PC12	PC13	PC14

Standard deviation	0.97046	0.86234	0.82617	0.70804	0.68580	0.62796	0.57121
Proportion of Variance	0.03139	0.02479	0.02275	0.01671	0.01568	0.01314	0.01088
Cumulative Proportion	0.84193	0.86672	0.88947	0.90618	0.92186	0.93501	0.94588
	PC15	PC16	PC17	PC18	PC19	PC20	PC21
Standard deviation	0.52326	0.48091	0.46611	0.45725	0.40475	0.3754	0.3332
Proportion of Variance	0.00913	0.00771	0.00724	0.00697	0.00546	0.0047	0.0037
Cumulative Proportion	0.95501	0.96272	0.96996	0.97693	0.98239	0.9871	0.9908
	PC22	PC23	PC24	PC25	PC26	PC27	PC28
Standard deviation	0.23327	0.2191	0.20233	0.19205	0.18114	0.16743	0.13202
Proportion of Variance	0.00181	0.0016	0.00136	0.00123	0.00109	0.00093	0.00058
Cumulative Proportion	0.99260	0.9942	0.99556	0.99679	0.99789	0.99882	0.99940
	PC29	PC30					
Standard deviation	0.10866	0.07828					
Proportion of Variance	0.00039	0.00020					
Cumulative Proportion	0.99980	1.00000					


Q4. From your results, what proportion of the original variance is captured by the first principal components (PC1)? PC1 captures approximately 38.96% of the original variance.

Q5. How many principal components (PCs) are required to describe at least 70% of the original variance in the data? 2 PCs are required to describe at least 70% of the original variance in the data.

Q6. How many principal components (PCs) are required to describe at least 90% of the original variance in the data? 7 PCs are required to describe at least 90% of the original variance in the data.

## 2. Interpreting PCA Results


```
biplot(wisc.pr)
```



Q7. What stands out to you about this plot? Is it easy or difficult to understand? Why? The center of the plot stands out to me only because it is so messy. It is difficult to understand given how cluttered it is.

```
# Scatter plot observations by components 1 and 2


plot(wisc.pr$x[, c(1, 2)],
      xlab = "PC1", ylab = "PC2")
```



Q8. Generate a similar plot for principal components 1 and 3. What do you notice about these plots?

```
plot(wisc.pr$x[, c(1, 3)],  
      xlab = "PC1", ylab = "PC3")
```






Answer: I was having issues with `col=diagnosis` so I omitted it in these plots (Error: unexpected symbol in: “`plot(wisc.pr$x[, c(1, 3)], col = diagnosis xlab)`” & invalid color name ‘diagnosis’) but I notice that PC2 has more variance than PC3.

## 2.Variance Explained

```
# Create a data.frame for ggplot
df <- as.data.frame(wisc.pr$x)
df$diagnosis <- diagnosis

# Load the ggplot2 package
library(ggplot2)

# Make a scatter plot colored by diagnosis
ggplot(df) +
  aes(PC1, PC2, col=diagnosis) +
  geom_point()
```




```
# Calculate variance of each component
pr.var <- wisc.pr$sdev^2
head(pr.var)
```

```
[1] 11.686620  3.441871  2.983241  1.979494  1.738785  1.359503
```


```
# Calculate total variance explained
total_var <- sum(pr.var)

# Variance explained by each principal component: pve
pve <- pr.var / total_var

# Plot variance explained for each principal component
plot(pve, xlab = "Principal Component",
     ylab = "Proportion of Variance Explained",
     ylim = c(0, 1), type = "o")
```



```
# Alternative scree plot of the same data, note data driven y-axis
barplot(pve, ylab = "Precent of Variance Explained",
        names.arg=paste0("PC",1:length(pve)), las=2, axes = FALSE)
axis(2, at=pve, labels=round(pve,2)*100 )
```



Q9. For the first principal component, what is the component of the loading vector (i.e. `wisc.pr$rotation[,1]`) for the feature `concave.points_mean`?

```
# Loading vector component for feature concave.points_mean
loading_component <- wisc.pr$rotation[, 1]
loading_component
```

X.1	X.2	X.3	X.4	X.5	X.6
-0.14492806	-0.11319439	0.11734576	0.02507453	-0.18373394	-0.27308944
X.7	X.8	X.9	X.10	X.11	X.12
-0.27475507	-0.25905963	-0.17075393	-0.12574144	-0.19310926	-0.05008037
X.13	X.14	X.15	X.16	X.17	X.18
-0.20612259	-0.10000261	-0.06523108	-0.23750442	-0.24022059	-0.23232505
X.19	X.20	X.21	X.22	X.23	X.24
-0.04820717	-0.19059524	-0.16802634	-0.11383556	0.13378332	0.09849681
X.25	X.26	X.27	X.28	X.29	X.30
-0.17625086	-0.25367150	-0.25621658	-0.25391455	-0.14386065	-0.19384792

Answer: X.8 is concave points mean, -0.25905963.

Q10. What is the minimum number of principal components required to explain 80% of the variance of the data?

```
# Cumulative proportion of variance explained
cumulative_pve <- cumsum(pve)

# Minimum number of principal components to explain 80% of the variance
min_components <- which.max(cumulative_pve >= 0.8)
```

Answer: 7

### 3. Hierarchical Clustering


```
# Scale the wisc.data data using the "scale()" function
data.scaled <- scale(wisc.data)

# Calculate the (Euclidean) distances between all pairs of observations in the new scaled
data.dist <- dist(data.scaled)

# Create a hierarchical clustering model using complete linkage
wisc.hclust <- hclust(data.dist, method = "complete")

# Plot the dendrogram
plot(wisc.hclust)
abline( col = "red", lty = 2)
```

## Cluster Dendrogram



```
data.dist
hclust(*, "complete")
```

Q11. Using the `plot()` and `abline()` functions, what is the height at which the clustering model has 4 clusters? Around 10 according to my plot. However the plot on the website is different and is around 19.

```
# Cut the tree into 4 clusters
wisc.hclust.clusters <- cutree(wisc.hclust, k = 4)

# Compare the cluster membership to the actual diagnoses.

table(wisc.hclust.clusters, diagnosis)
```

	diagnosis			
wisc.hclust.clusters	B	diagnosis	M	
1	11	1	114	
2	3	0	53	
3	173	0	40	
4	170	0	5	

Q12. Can you find a better cluster vs diagnoses match by cutting into a different number of clusters between 2 and 10?

```
# Cut the tree into 10 clusters
wisc.hclust.clusters <- cutree(wisc.hclust, k = 2)

# Compare the cluster membership to the actual diagnoses.

table(wisc.hclust.clusters, diagnosis)
```

	diagnosis		
wisc.hclust.clusters	B	diagnosis	M
1	11	1	114
2	346	0	98

Q13. Which method gives your favorite results for the same data.dist dataset? Explain your reasoning.

```
# Create hierarchical clustering models using different methods
wisc.hclust_single <- hclust(data.dist, method = "single")
wisc.hclust_complete <- hclust(data.dist, method = "complete")
wisc.hclust_average <- hclust(data.dist, method = "average")
wisc.hclust_ward <- hclust(data.dist, method = "ward.D2")

# Cut the trees into clusters (e.g., let's use 4 clusters for comparison)
wisc.hclust_single_clusters <- cutree(wisc.hclust_single, k = 4)
wisc.hclust_complete_clusters <- cutree(wisc.hclust_complete, k = 4)
wisc.hclust_average_clusters <- cutree(wisc.hclust_average, k = 4)
wisc.hclust_ward_clusters <- cutree(wisc.hclust_ward, k = 4)

# Compare cluster vs. diagnosis match for each method
print("Single linkage:")
```

```
[1] "Single linkage:"
```

```
print(table(wisc.hclust_single_clusters, diagnosis))
```

	diagnosis		
wisc.hclust_single_clusters	B	diagnosis	M
1	0	1	0
2	355	0	211
3	0	0	1
4	2	0	0

```
print("\nComplete linkage:")
```

```
[1] "\nComplete linkage:"
```

```
print(table(wisc.hclust_complete_clusters, diagnosis))
```

	diagnosis			
wisc.hclust_complete_clusters	B	diagnosis	M	
1	11		1	114
2	3		0	53
3	173		0	40
4	170		0	5

```
print("\nAverage linkage:")
```

```
[1] "\nAverage linkage:"
```

```
print(table(wisc.hclust_average_clusters, diagnosis))
```

	diagnosis			
wisc.hclust_average_clusters	B	diagnosis	M	
1	0		1	0
2	19		0	182
3	331		0	29
4	7		0	1

```
print("\nWard linkage:")
```

```
[1] "\nWard linkage:"
```

```
print(table(wisc.hclust_ward_clusters, diagnosis))
```



		diagnosis		
wisc.hclust_ward_clusters	B	diagnosis	M	
1	1	1	177	
2	58	0	30	
3	165	0	3	
4	133	0	2	

I like the single clusters method most because the numbers are easier to look at.

## 5. Combining Methods


```
# Find the minimum number of principal components required to describe at least 90% of the
cumulative_variance <- cumsum(wisc.pr$sdev^2 / sum(wisc.pr$sdev^2))
num_components_90 <- which(cumulative_variance >= 0.9)[1]

# Use the first num_components_90 principal components
wisc.pr_reduced <- wisc.pr$x[, 1:num_components_90]

# Create hierarchical clustering model with ward.D2 linkage method
wisc.pr.hclust <- hclust(dist(wisc.pr_reduced), method = "ward.D2")

plot(wisc.pr.hclust)
```

## Cluster Dendrogram



```
dist(wisc.pr_reduced)
hclust (*, "ward.D2")
```


```
grps <- cutree(wisc.pr.hclust, k=2)
table(grps)
```

```
grps
  1   2
352 218
```


```
table(grps, diagnosis)
```

```
diagnosis
grps  B diagnosis  M
  1 151          1 200
  2 206          0  12
```

```
plot(wisc.pr$x[,1:2], col=grps)
```



```
plot(wisc.pr$x[,1:2])
```




```
g <- as.factor(grps)
levels(g)
```

```
[1] "1" "2"
```

```
g <- relevel(g,2)
levels(g)
```

```
[1] "2" "1"
```

```
# Plot using our re-ordered factor
plot(wisc.pr$x[,1:2], col=g)
```



```
# Use the distance along the first 7 PCs for clustering
wisc.pr.hclust <- hclust(dist(wisc.pr$x[, 1:7]), method = "ward.D2")

# Cut the hierarchical clustering model into 2 clusters
wisc.pr.hclust.clusters <- cutree(wisc.pr.hclust, k = 2)
```

Q15. How well does the newly created model with four clusters separate out the two diagnoses?

```
# Compare the results from the new hierarchical clustering model with the actual diagnoses
table(wisc.pr.hclust.clusters, diagnosis)
```

	diagnosis		
wisc.pr.hclust.clusters	B	diagnosis	M
1	16	1	185
2	341	0	27

The diagnoses are separated better in this model.

Q16. How well do the k-means and hierarchical clustering models you created in previous sections (i.e. before PCA) do in terms of separating the diagnoses? Again, use the `table()` function to compare the output of each model (`wisc.km$cluster` and `wisc.hclust.clusters`) with the vector containing the actual diagnoses.

```
hclust_diagnosis_table <- table(wisc.hclust.clusters, diagnosis)
```

“Error: object ‘wisc.km’ not found.”

## 6. Sensitivity/Specificity

```
# Function to calculate specificity
calculate_specificity <- function(confusion_matrix) {
  true_negatives <- confusion_matrix[1, 1]
  false_positives <- confusion_matrix[1, 2]
  return(true_negatives / (true_negatives + false_positives))
}

# Function to calculate sensitivity
calculate_sensitivity <- function(confusion_matrix) {
  true_positives <- confusion_matrix[2, 2]
  false_negatives <- confusion_matrix[2, 1]
  return(true_positives / (true_positives + false_negatives))
}

# Calculate specificity and sensitivity for each clustering model
```

```

hclust_specificity <- calculate_specificity(hclust_diagnosis_table)
hclust_sensitivity <- calculate_sensitivity(hclust_diagnosis_table)

# Display the results
print(paste("Hierarchical clustering specificity:", hclust_specificity))

```

```
[1] "Hierarchical clustering specificity: 0.916666666666667"
```

```

print(paste("Hierarchical clustering sensitivity:", hclust_sensitivity))

```

```
[1] "Hierarchical clustering sensitivity: 0"
```

I keep getting errors and cannot figure them out. But based on what is on the website, I would say the second model is better at separating the diagnoses.

Q17. Which of your analysis procedures resulted in a clustering model with the best specificity? How about sensitivity? The second model is more specific and the second is more sensitive.


## 7. Prediction

```

url <- "new_samples.csv"
url <- "https://tinyurl.com/new-samples-CSV"
new <- read.csv(url)
npc <- predict(wisc.pr)

plot(wisc.pr$x[,1:2], col=g)
points(npc[,1], npc[,2], col="blue", pch=16, cex=3)
text(npc[,1], npc[,2], c(1,2), col="white")

```



Q18. Which of these new patients should we prioritize for follow up based on your results?  
 My plot came out too messy. Based on one on website I would say patients 1 and 2.